

Tree based methods I

Regression tree example

STAT 32950-24620

Spring 2023 (5/16)

1 / 24

Tree based methods

- Supervised learning, with explanatory variables and a response variable.
- Find a partition of the space of explanatory variables (predictors, features).
- Provide a constant prediction in each partition region.

2 / 24

Tree based methods

- **Regression Tree** for estimating a continuous response variable
 - Take an average as the prediction in each partition region.
- **Classification Tree** for classifying a categorical response var.
 - Take a vote as the prediction in each partition region.
- Partition of the explanatory variable space is obtained by splitting the range of one predictor at a time.
- The partition regions are rectangles or hyper-rectangles.

3 / 24

Regression Tree

Regression Tree: the response variable is continuous.
(vs Classification Tree: the response variable is categorical)

```
# libraries for Regression Tree
library(MASS)
library(tree)      # earlier than "rpart" package
library(rpart)     # alternative to "tree" package
library(rpart.plot) # nicer tree plots
```

rpart — Recursive Partitioning And Regression Trees

4 / 24

Data Example

Boston housing data

$n = 506$ observations, $p = 13$ input variables.

Response variable:

medv: median value of owner-occupied homes in USD 1000's

5 / 24

Example data: Explanatory variables

- crim: per capita crime rate by town
- zn: % residential land zoned for lots over 25k sq.ft
- indus: % of non-retail business acres per town
- chas: Charles River dummy variable
(= 1 if tract bounds river; 0 otherwise)
- nox: nitric oxides concentration (parts per 10 million)
- rm: average number of rooms per dwelling
- age: % of owner-occupied units built prior to 1940
- dis: weighted distances to 5 Boston employment centres
- rad: index of accessibility to radial highways
- tax: full-value property-tax rate per USD 10,000
- ptratio: pupil-teacher ratio by town
- black: $1000(B_k - 0.63)^2$, B_k = % of blacks by town
- lstat: percentage of lower status of the population

6 / 24

Check data format

```
str(Boston); #summary(Boston)
```

```
## 'data.frame': 506 obs. of 14 variables:
## $ crim : num 0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn : num 18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus : num 2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 ...
## $ chas : int 0 0 0 0 0 0 0 0 0 0 ...
## $ nox : num 0.538 0.469 0.469 0.458 0.458 0.458 0.528 0.477 ...
## $ rm : num 6.58 6.42 7.18 7 7.15 ...
## $ age : num 65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 ...
## $ dis : num 4.09 4.97 4.97 6.06 6.06 ...
## $ rad : int 1 2 2 3 3 3 5 5 5 5 ...
## $ tax : num 296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 ...
## $ black : num 397 397 393 395 397 ...
## $ lstat : num 4.98 9.14 4.03 2.94 5.33 ...
## $ medv : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 18.5 22.0 ...
```

7 / 24

Regression Tree on the training data

For this dataset, we split the data into two equal sizes,
create training and testing data.

```
set.seed(1)
# random sample from row numbers
train = sample(1:nrow(Boston), nrow(Boston)/2)
```

Fit a regression tree on the training data:

```
tree.boston = tree(medv ~ ., Boston, subset = train)
```

8 / 24

Check the fitted tree


```
summary(tree.boston)
```

```
##
## Regression tree:
## tree(formula = medv ~ ., data = Boston, subset = train)
## Variables actually used in tree construction:
## [1] "rm" "lstat" "crim" "age"
## Number of terminal nodes: 7
## Residual mean deviance: 10.4 = 2550 / 246
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -10.200  -1.780  -0.177   0.000   1.920   16.600
```

9 / 24

Plot the fitted regression tree (using tree)

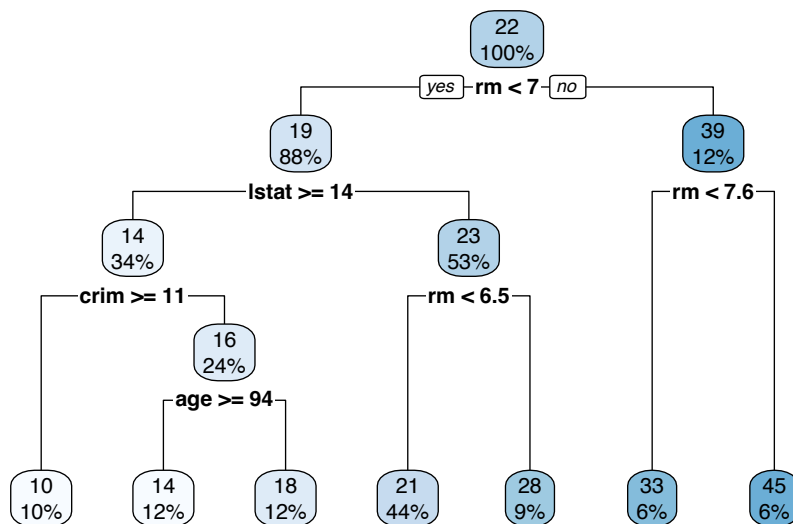
```
plot(tree.boston)
text(tree.boston, cex=0.7, digits=2)
```



10 / 24

Plot the regression tree (using rpart)

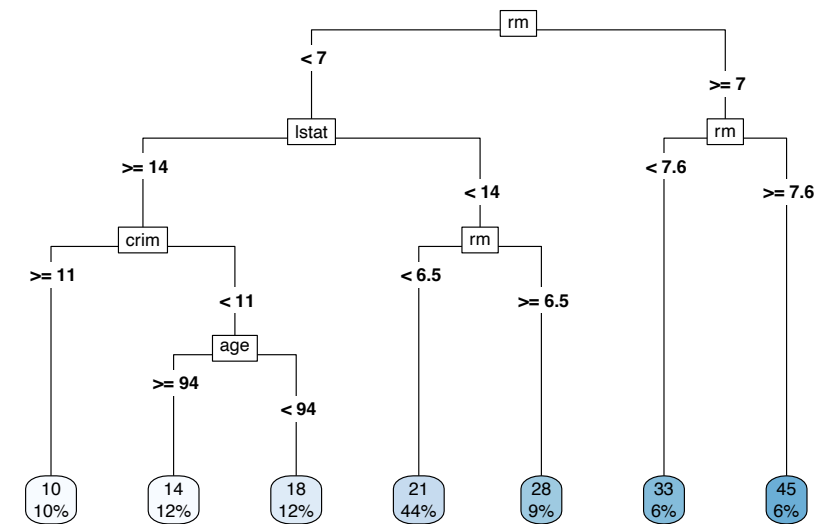
```
rpart.plot(rpart(medv~., Boston, subset=train))
```



11 / 24

Plot the tree in another style (using rpart)


```
rpart.plot(rpart(medv~., Boston, subset=train), type=5)
```



12 / 24

Feature space stratification: Step 1 of the split


```
plot(Boston[train,]$rm,Boston[train,]$medv,cex=.5,  
      xlab="rm",ylab="medv"); abline(v=7, col=2)
```



13 / 24

Feature space stratification: Steps 2


```
plot(Boston[train,]$rm,Boston[train,]$lstat,cex=.5,  
      xlab="rm",ylab="lstat"); abline(v=7,col=2); segments(3,14.4, 7,14.4,col=3)
```



14 / 24

Feature space stratification: Steps 3

```
plot(Boston[train,]$rm,Boston[train,]$lstat,cex=.5,  
      xlab="rm",ylab="lstat"); abline(v=7,col=2);  
segments(3,14.4, 7,14.4,col=3); abline(v=7.6,col=4)
```



15 / 24

How to build a regression tree

- Grow the tree upside down from the root to leaves.
- Start with a single region R_k , iterate.
 - Select a region R_k , a predictor X_j , a splitting point s , such that splitting R_k with the rule $X_j < s$ optimally reduce RSS


$$\sum_{m=1}^{|T|} \sum_{x_i \in R_m} (y_i - \bar{y}_{R_m})^2$$

- Redefine the regions with the additional split.
- Iterate.
- Terminate when there are very few observations (e.g. ≤ 5) in a region.

16 / 24

Tree deviance vs tree size


```
cv.boston=cv.tree(tree.boston)
plot(cv.boston$size,cv.boston$dev,type='b')
```



17 / 24

Prune tree using deviance (5 regions)


```
plot(prune.tree(tree.boston,best =5)) # use best=tree-size
text(prune.tree(tree.boston,best =5),cex=0.7,digits =2)
```



18 / 24

Prune tree using deviance (3 regions)


```
plot(prune.tree(tree.boston,best =3))
text(prune.tree(tree.boston,best =3))
```



19 / 24

Tree fit


```
yhat=predict(tree.boston,newdata=Boston[-train,])
boston.test=Boston[-train,"medv"]
plot(yhat,boston.test,cex=.5); abline(0,1)
```



20 / 24

Linear regression fit


```
yhatReg=predict(lm(medv~.,Boston,subset=train),
               newdata=Boston[-train,])
plot(yhatReg,Boston[-train,"medv"],cex=.5); abline(0,1)
```



21 / 24

Comparison of prediction of tree vs lm

```
yhat=predict(tree.boston,newdata=Boston[-train,])
boston.test=Boston[-train,"medv"]
plot(boston.test,yhat,cex=.5); abline(0,1)
points(Boston[-train,"medv"],yhatReg,cex=.5,pch=4,col=2)
```



22 / 24

Comparison of mean SS Residuals of tree vs lm

```
mean((yhat-boston.test)^2)
## [1] 35.29
mean((yhatReg-boston.test)^2)
## [1] 26.86
```

Which method is better?

Note: comparison of mean SS residual of training and testing data.

23 / 24

Cross validation

- Split the training data into 10 folds (default)
- For $k = 1, \dots, 10$, using every fold except the k th
- Construct trees T_1, \dots, T_m for a range of α in

$$\min_T \left(\sum_{m=1}^{|T|} \sum_{x_i \in R_m} (y_i - \bar{y}_{R_m})^2 + \alpha |T| \right)$$

Cost complexity pruning

- For each tree T_i , calculate RSS on the test set
- Remove the **Weakest link**, the subtree minimizes

$$\frac{RSS(T_1) - RSS(T_0)}{|T_0| - |T_1|}$$

- Select α that minimizes the average testing error.

24 / 24